

▼ Example of Linear Regression with Two Input Variables using the Iris Flower Dataset

- We would like to predict "petal width" (column 4 in the original dataset) using sepal width (column 2) and petal length (column 3)
- Metadata: <https://github.com/badriadhikari/2019-Spring-AI/blob/master/supplementary/iris.names>

```
1 from keras.models import Sequential
2 from keras.layers import Dense
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 # Column 2. sepal width in cm (load as col 0)
7 # Column 3. petal length in cm (load as col 1)
8 # Column 4. petal width in cm (load as col 2)
9 datapath = "https://raw.githubusercontent.com/badriadhikari/2019-Spring-AI/"
10 datapath = datapath + "master/supplementary/iris.data"
11 dataset = np.genfromtxt(datapath, delimiter=",", usecols=(1, 2, 3))
12
13 print('')
14 print(dataset.shape)
15 print('')
16 print(dataset[0:5])
```

☞ Using TensorFlow backend.

```
(150, 3)
```

```
[[3.5 1.4 0.2]
 [3.  1.4 0.2]
 [3.2 1.3 0.2]
 [3.1 1.5 0.2]
 [3.6 1.4 0.2]]
```

```
1 # Q1. Why is shuffling important before splitting?
2 np.random.shuffle(dataset)
3 print('')
4 print(dataset[0:5])
5 train = dataset[:100]
6 valid = dataset[100:]
7 print('')
8 print(train.shape)
9 print('')
10 print(valid.shape)
```

☞

```
[[2.7 5.3 1.9]
 [3.2 6.  1.8]
 [3.  5.1 1.8]
 [2.4 3.8 1.1]
 [3.  4.5 1.5]]
```

```
(100, 3)
```


```
(50, 3)
```

```
1 #Q2. Which of the two input features seems more useful
2 #    for predicting petal width?
```

```

3 plt.figure(figsize=(4,4))
4 plt.scatter(train[:, 0], train[:, 2], color = 'r', alpha = 0.5)
5 plt.xlabel('sepal width in cm')
6 plt.ylabel('petal width in cm')
7 plt.show()
8 plt.figure(figsize=(4,4))
9 plt.scatter(train[:, 1], train[:, 2], color = 'b', alpha = 0.5)
10 plt.xlabel('petal length in cm')
11 plt.ylabel('petal width in cm')
12 plt.show()

```



```

1 train_input = train[:, 0:2] # col 2 & 3
2 train_output = train[:, 2] # col 4
3 valid_input = valid[:, 0:2]
4 valid_output = valid[:, 2]
5
6 print('')
7 print(train_input[0:5])
8 print('')
9 print(train_output[0:5])

```



```
rr2.7 5.31
```


```
1 #Q3. Why is the number of parameters = 3?
2 model = Sequential()
3 model.add(Dense(1, input_dim = len(train_input[0]), activation='linear'))
4 print(model.summary())
5
6 # Changing 'mae' to 'mse' should improve the smoothness of
7 # the learning curve and possibly the overall errors
8 model.compile(loss='mae', optimizer='sgd', metrics=['mae'])
9
10 # Verbose = 0 shows no updates, can be changed to 1 or 2
11 history = model.fit(train_input, train_output, epochs=50,
12                     verbose = 0, batch_size=10,
13                     validation_data = (valid_input, valid_output))
14
```



Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 1)	3

=====
Total params: 3
Trainable params: 3
Non-trainable params: 0
=====
None

```
1 #Q4. Why eventually validation MAE is not
2 # always less than train MAE?
3 plt.figure(figsize=(4,4))
4 plt.plot(history.history['mean_absolute_error'])
5 plt.plot(history.history['val_mean_absolute_error'])
6 plt.ylabel('mean absolute error')
7 plt.xlabel('epoch')
8 plt.legend(['train', 'validation'], loc='upper right')
9 plt.show()
```



```
1 #Q5. Are these predictions reasonable?
2 np.set_printoptions(precision = 2)
3 print ('True Validation Data:')
4 print(valid_output[0:5])
5 prediction = model.predict(valid_input)
```